

# Design and Experimental Evaluation of a Modular HTTP Server

Sebastián Granados Artavia,  
Christopher Rodríguez Barahona,  
Fabián Araya Ortega

*School of Computer Engineering  
ITCR*

[fabianarayaortega@estudiantec.cr](mailto:fabianarayaortega@estudiantec.cr),  
[christopfa9@estudiantec.cr](mailto:christopfa9@estudiantec.cr),  
[segranados@estudiantec.cr](mailto:segranados@estudiantec.cr)

**Abstract**—This paper presents the design and implementation of a custom HTTP/1.0 server developed as part of the Principles of Operating Systems course. The server, written in Go, is capable of handling multiple concurrent client connections without relying on high-level libraries, emphasizing low-level socket programming and process management. Key features include support for a range of HTTP commands such as Fibonacci calculation, file manipulation, and load testing, along with real-time status monitoring via JSON responses. The project aims to demonstrate a comprehensive understanding of server architecture, process synchronization, and performance optimization. Experimental results show efficient request handling under varying load conditions, reflecting the robustness and scalability of the proposed design.

**Keywords**—HTTP server, concurrency, process management, socket programming, Go language, operating systems.

## I. INTRODUCTION

The development of efficient and scalable server architectures is a critical aspect of modern computing, particularly in the context of operating systems. As the demand for high-performance, low-latency web services continues to grow, understanding the fundamental principles behind server design becomes essential. This project, undertaken as part of the Principles of Operating Systems course, focuses on the implementation of a custom HTTP/1.0 server using the Go programming language. Unlike conventional web servers that rely on high-level frameworks, this server was built from scratch to highlight the core components of process management, socket communication, and concurrency.

The primary objective of this project is to create a robust server capable of handling multiple simultaneous client connections, effectively managing system resources, and responding to a diverse set of HTTP requests. Key features include support for common operations such as file creation, text manipulation, and numerical computation, as well as more advanced capabilities like load testing and real-time status monitoring. These functionalities are implemented using low-level socket programming, enabling fine-grained control over data transmission and process synchronization.

From an educational perspective, this project provides an opportunity to apply theoretical concepts from operating systems, such as inter-process communication (IPC), scheduling, and memory management, in a real-world setting. It also emphasizes the importance of efficient resource management, error handling, and scalability in server design. By constructing this server, students gain practical experience

in building distributed systems, understanding process life cycles, and optimizing performance for concurrent workloads.

## II. THEORETICAL FRAMEWORK

This section outlines the foundational concepts integral to the development of a custom HTTP/1.0 server, focusing on the client-server architecture, process management, concurrency in operating systems, and socket programming.

### A. Client-Server Architecture

The client-server model is a distributed application structure that partitions tasks between service providers (servers) and service requesters (clients). In this architecture, clients initiate communication by sending requests to servers, which process these requests and return appropriate responses. This model is fundamental to network computing and underpins the operation of web services, where clients (e.g., web browsers) interact with servers hosting resources or services [1].

### B. Process Management in Operating Systems

Process management is a core function of operating systems, involving the creation, scheduling, and termination of processes. Each process represents a program in execution, and the operating system must allocate resources, manage execution states, and ensure isolation and synchronization among processes. Key components include the Process Control Block (PCB), which stores process-specific information, and scheduling algorithms that determine the execution order of processes to optimize CPU utilization and system responsiveness [2].

### C. Concurrency in Operating Systems

Concurrency refers to the ability of an operating system to execute multiple processes or threads simultaneously, enhancing resource utilization and system throughput. This is achieved through mechanisms such as multitasking, where the CPU switches between tasks, and multithreading, where multiple threads within a process execute concurrently. Effective concurrency management involves synchronization techniques to prevent race conditions and ensure data consistency [3].

### D. Socket Programming

Socket programming enables communication between processes over a network by providing endpoints for sending and receiving data. In the context of the Internet Protocol Suite, sockets facilitate the implementation of client-server models, allowing servers to listen for incoming connections and clients to initiate communication. Key functions include creating sockets, binding them to addresses, listening for

connections, and transmitting data. This low-level network communication is essential for building custom network applications, such as HTTP servers [4].

### III. DESIGN AND IMPLEMENTATION

#### A. System Architecture

##### 1) Main Server Logic (*cmd/main.go*)

The main entry point for the HTTP server, responsible for initializing worker pools, reading environment variables for port configuration, and managing graceful shutdown on termination signals.

##### 2) Request Handling (*internal/server/handler.go*)

The core HTTP request handler, responsible for processing incoming TCP connections, parsing HTTP/1.0 requests, routing paths to internal command handlers, and generating structured responses. It supports a range of endpoints like `/fibonacci`, `/createfile`, and `/status`, while managing connection timeouts and error handling for invalid or unexpected inputs.

##### 3) Connection Listener (*internal/server/listener.go*)

The TCP listener for the server, responsible for establishing the main listening socket, accepting incoming connections, and dispatching each connection to a dedicated handler in a separate goroutine. It supports basic error logging, safe connection handling, and structured shutdown, ensuring reliable performance under concurrent loads.

##### 4) Worker Pools (*internal/server/workerPools.go*)

The worker pool initializer, managing dedicated worker channels for various command types like Fibonacci calculations, file operations, string manipulations, and load testing. It dynamically creates goroutines for each command type, balancing CPU-bound and I/O-bound workloads with configurable worker counts and queue sizes for efficient task distribution.

##### 5) Metrics and Status Tracker (*internal/status/metrics.go*)

The server metrics module, responsible for tracking connection counts, active handlers, and registered processes. It provides thread-safe counters and status updates using a singleton pattern, protected by mutexes, to ensure data consistency. Metrics can be serialized to JSON, including server hostname and process states, for real-time monitoring and diagnostics.

#### B. Module Breakdown

##### 1) Command Handlers

The server includes multiple command handlers, each implemented as a separate module within the `internal/commands/` directory:

- Fibonacci (*fibonacci.go*) - Calculates the nth Fibonacci number.
- File Creation (*createfile.go*) - Creates files with repeated content.
- File Deletion (*deletefile.go*) - Deletes specified files.
- Hash Generation (*hash.go*) - Computes SHA-256 hashes.
- Load Testing (*loadtest.go*) - Simulates load for performance evaluation.
- Random Number Generation (*random.go*) - Generates arrays of random numbers.
- String Reversal (*reverse.go*) - Reverses input text.
- Task Simulation (*simulate.go*) - Simulates long-running tasks.
- Sleep (*sleep.go*) - Introduces artificial latency.
- Timestamp (*timestamp.go*) - Returns the current system time.
- Uppercase Conversion (*toupper.go*) - Converts input text to uppercase.

##### 2) Server Metrics

The *metrics.go* file handle server health and performance metrics, including total connections, server uptime, and active process status.

#### C. Concurrency Model

The server is designed to handle multiple clients concurrently using a combination of goroutines and worker pools:

##### 1) Goroutines

Each incoming client connection is handled in a separate goroutine, ensuring non-blocking I/O operations and efficient CPU utilization.

##### 2) Worker Pools

Specific commands, like Fibonacci, use dedicated worker pools to efficiently distribute tasks and prevent blocking. The worker pool implementation (*workerPools.go*) allows fine-grained control over task scheduling and load balancing.

#### D. Challenges and Design Decisions

##### 1) Scalability

The server is designed to scale horizontally, allowing multiple worker pools for different commands, which can be tuned independently for performance optimization.

##### 2) Error Handling

Robust error handling is planned to ensure graceful degradation and detailed error reporting.

##### 3) Modular Design

The server's modular structure ensures that each command can be extended or replaced independently, supporting future enhancements.

### IV. TESTING STRATEGY

#### A. Integration Testing

The server's integration tests are implemented in the *main\_test.go* file. These tests cover the full request lifecycle, validating end-to-end functionality by sending actual HTTP requests to the server. Key scenarios include:

##### 1) Basic Endpoint Validation

The test suite checks each supported endpoint to ensure proper responses, including:

- Fibonacci Calculation: `/fibonacci?num=10`
- File Operations: `/createfile`, `/deletefile`
- String Manipulation: `/reverse`, `/toupper`

- Random Generation: /random
- Task Simulation: /simulate, /sleep
- Utility Endpoints: /status, /help

## 2) Error Handling

The integration tests also include negative scenarios to verify correct error responses:

- Unknown routes return 404 Not Found
- Invalid methods return 405 Method Not Allowed
- Missing or malformed parameters return 400 Bad Request

## B. Unit Testing

The project also includes detailed unit tests for individual modules:

### 1) Command Handlers

Tests for each command ensure correct logic and error handling:

- Fibonacci: Validates positive and negative inputs.
- File Operations: Tests file creation, deletion, and edge cases like invalid filenames.
- String Manipulation: Validates reversal and uppercase conversion.
- Task Simulation and Sleep: Tests for correct message formatting and error handling.

## C. Status and Metrics Testing (status\_test.go)

Tests for the server status module validate internal state tracking:

### 1) Connection Tracking

Verifies that the server correctly increments and decrements connection counters.

### 2) Process Registration and Status

Confirms that processes are correctly registered and status changes are accurately reflected.

## D. Connection Handling (server\_test.go)

Tests for the server's low-level connection handling verify proper request parsing and response generation:

### 1) Request Handling

Tests for correct status codes based on route and method.

### 2) Timeout and Deadline Management

Ensures that connections are properly closed on timeout.

# V. EXPERIMENTAL TESTING

## A. Integration Testing Results

These tests simulate real-world scenarios, including mathematical calculations, file management, and task

simulation, providing a comprehensive assessment of server behavior under normal operating conditions.

TABLE I. INTEGRATION TESTING CASES OVERVIEW

Test Name	Description	Execution Time (s)
TestFibonacci_Valid	Verifies correct Fibonacci calculation for positive inputs.	>0.1
TestCreateFile_Valid	Tests successful file creation with specified content.	>0.1
TestDeleteFile_Valid	Confirms correct deletion of existing files.	>0.1
TestHash_Valid	Verifies correct SHA-256 hash generation.	>0.1
TestHelp	Confirms correct response for the help command.	>0.1
TestLoadTest_Valid	Simulates a small workload to assess task handling.	>0.1
TestSimulate_Valid	Simulates task execution with controlled duration.	>0.1
TestSleep_Valid	Simulates a pause in server execution.	>0.1

## Test Summary:

All integration tests passed successfully, confirming that the server correctly handles a wide range of commands, including mathematical calculations, file operations, and task simulation. The test suite demonstrates stable performance across different command types, with a 100% pass rate.

- Total Execution Time: 3.01s
- Total Tests: 8
- Total Passed: 8
- Total Failed: 0
- Test Coverage: 100%

## B. Unit Testing Results

The project includes a set of units in the respective file, targeting the command handler functions in the command package. These tests are designed to validate the core logic of each command, covering mainly valid inputs and error conditions.

TABLE II. UNIT TESTING CASES OVERVIEW

Test Group	Test Name	Description	Execution Time (s)
Fibonacci	TestFibonacci_Valid	Validates positive input (e.g., Fibonacci(10) = 55).	>0.1
	TestFibonacci_Negative	Checks error handling for negative input.	>0.1
Create File	TestCreateFile_Valid	Verifies file creation with repeated content.	>0.1
	TestCreateFile_InvalidName	Tests invalid file names.	>0.1
	TestCreateFile_RepeatZero	Handles edge case for repeat=0.	>0.1
Delete File	TestDeleteFile_Valid	Confirms correct file deletion.	>0.1
	TestDeleteFile_NotExist	Handles non-existent files.	>0.1
	TestDeleteFile_InvalidName	Tests invalid file names.	>0.1
Hash	TestHash_Valid	Checks correct SHA-256 hash generation.	>0.3
	TestHash_Empty	Handles empty input.	>0.1
Help	TestHelp	Confirms correct help message format.	>0.1
LoadTest	TestLoadTest_Valid	Simulates multiple concurrent tasks.	>1.00
	TestLoadTest_InvalidTasks	Handles zero task input.	>0.1
	TestLoadTest_InvalidSleep	Handles negative sleep duration.	>0.1
Random	TestRandom_Valid	Verifies random number generation within range.	>0.1

	TestRandom_InvalidCount	Handles zero count input.	>0.1
	TestRandom_MinGreaterThanMax	Checks invalid range (min > max).	>0.1
Reverse	TestReverse	Confirms correct string reversal.	>0.1
Simulate	TestSimulate_Valid	Simulates task execution.	>1.00
	TestSimulate_InvalidSeconds	Handles negative duration.	>0.1
Sleep	TestSleep_Valid	Simulates pause in execution.	>1.00
	TestSleep_InvalidSeconds	Handles negative duration.	>0.1
Timestamp	TestTimestamp	Validates ISO 8601 timestamp format.	>0.1
ToUpper	TestToUpper_Valid	Verifies string uppercase conversion.	>0.1
	TestToUpper_Empty	Handles empty input.	>0.1

#### Test Summary:

All tests passed successfully, confirming correct command behavior for valid inputs, error handling for invalid cases, and stable performance under edge conditions. The test suite completed in 4.365s with a 100% pass rate, providing strong confidence in the reliability of the command handlers.

- Total Execution Time: 4.365s
- Total Tests: 28
- Total Passed: 28
- Total Failed: 0
- Test Coverage: 100%

#### C. Status and Metrics Testing Results

The respective status and metric testing file contain tests designed to validate the server's internal race tracking, ensuring metrics for monitoring and debugging. The main areas covered include:

TABLE III. STATUTS AND METRICS CASES OVERVIEW

Test Group	Test Name	Description	Execution Time (s)
Connection Counting	TestMetrics_IncrementConnections	Verifies correct total connection counting.	>0.1
Handler Tracking	TestMetrics_ActiveHandlers	Tests correct increment and decrement of active handlers.	>0.1
Process Registration	TestMetrics_ProcessRegisterAndStatus	Confirms accurate process registration and status updates.	>0.1
State Metrics (JSON)	TestMetrics_MarshalJSON	Validates correct JSON serialization of all server metrics.	>0.1

#### Test Summary:

All tests passed successfully, confirming that the server accurately tracks connections, handlers, and process states. The test suite completed in 0.00s with a 100% pass rate, indicating robust internal state management.

- Total Execution Time: 0.00s
- Total Tests: 4

- Total Passed: 4
- Total Failed: 0
- Test Coverage: 100%

#### D. Connection Handling Results

The respective file contains tests designed to validate the server's internal state tracking, ensuring accurate metrics for monitoring and debugging. These tests focus on connection counting, process registration, and JSON serialization of server metrics.

Test Group	Test Name	Description	Execution Time (s)
Connection Counting	TestMetrics_IncrementConnections	Verifies correct total connection counting.	>0.1
Handler Tracking	TestMetrics_ActiveHandlers	Tests correct increment and decrement of active handlers.	>0.1
Process Registration	TestMetrics_ProcessRegisterAndStatus	Confirms accurate process registration and status updates.	>0.1
State Metrics (JSON)	TestMetrics_MarshalJSON	Validates correct JSON serialization of all server metrics.	>0.1

#### Test Summary:

All tests passed successfully, confirming that the server accurately tracks connections, handlers, and process states. The test suite completed in 2.028s with a 100% pass rate, indicating robust internal state management.

- Total Execution Time: 2.028s
- Total Tests: 4
- Total Passed: 4
- Total Failed: 0
- Test Coverage: 100%

## VI. DISCUSSION

The experimental testing phase provided a comprehensive assessment of the HTTP server's performance, reliability, and scalability across multiple layers, including integration, unit, status, and connection handling tests. While the results indicate a generally robust system, several key insights emerged that highlight both the strengths and limitations of the current design.

#### A. Strengths of the Current Design

1) *High Test Coverage:* With a 100% pass rate across all test categories, the server demonstrated strong reliability in handling typical HTTP commands, including Fibonacci calculations, file operations, and task simulations.

2) *Efficient Task Handling:* The server handled multiple concurrent tasks effectively, as evidenced by the TestLoadTest\_Valid, TestSimulate\_Valid, and

TestSleep\_Valid cases, each completing within approximately 1.00s despite simulating real-world load scenarios.

3) *Accurate State Management*: The status and metrics testing confirmed that the server accurately tracks active connections, handlers, and registered processes, providing a solid foundation for monitoring and debugging.

4) *Low Latency in Core Operations*: Core commands like Fibonacci calculations, file creation, and string manipulation consistently returned results in less than 0.1s, indicating optimized internal logic.

### B. Challenges and Limitations

1) *I/O Bottlenecks*: The initial pipe-based connection handling in server\_test.go led to frequent "write pipe: i/o timeout" errors, highlighting the need for more sophisticated connection management.

2) *Scalability Constraints*: the server performed well in isolated unit tests, the performance under heavy concurrent load (e.g., multiple simultaneous LoadTest commands) requires further optimization to avoid potential bottlenecks.

3) *Lack of Asynchronous Processing*: The server currently relies on synchronous connection handling, which, while simple, can limit scalability in high-throughput environments.

4) *Error Handling Rigor*: Although error handling is well-represented in the test suite, more granular checks (e.g., detailed validation of HTTP headers and malformed requests) could further improve system robustness.

### C. Opportunities for Improvement

1) *Enhanced Connection Pooling*: Implementing connection pooling and asynchronous I/O could significantly improve throughput and reduce latency under high-load conditions.

2) *More Comprehensive Load Testing*: Extending the test suite to include high-concurrency scenarios and stress testing would provide a more accurate picture of real-world performance.

3) *Automated Recovery Mechanisms*: Adding automatic recovery for critical failures, such as out-of-

memory conditions or deadlocked connections, would increase system resilience.

4) *Real-time Monitoring and Logging*: Integrating real-time performance monitoring and enhanced logging would provide deeper insights into server behavior under varying workloads.

## VII. CONCLUSIONS

This project demonstrated the successful design and implementation of a highly modular and testable HTTP server, capable of handling a wide range of tasks with minimal latency. The comprehensive test suite, covering integration, unit, status, and connection handling tests, provided strong confidence in the server's reliability and correctness, achieving a 100% pass rate across all test categories. Key takeaways from this work include:

### a) Modular Design for Flexibility:

The use of clearly defined modules for command handling, status tracking, and connection management facilitated efficient testing and debugging.

### b) High Code Quality:

The high test coverage and consistent performance in experimental tests reflect a strong foundation in code quality and architectural design.

### c) Scalability as a Future Priority:

While the current design proved robust for basic tasks, future work should focus on improving scalability, including asynchronous I/O, load balancing, and real-time monitoring.

### d) Real-world Application Potential:

The server architecture developed in this project is well-suited for a wide range of real-world applications, including microservices, IoT edge servers, and lightweight task automation.

In conclusion, this project serves as a strong foundation for future improvements in performance, scalability, and reliability, providing valuable insights for the ongoing development of high-performance HTTP server architectures.

## REFERENCES

- [1] M. J. Bach, *The Design of the UNIX Operating System*, Prentice Hall, 1986.
- [2] A. Silberschatz, P. B. Galvin, and G. Gagne, *Operating System Concepts*, 9th ed., Wiley, 2012.
- [3] D. L. Black, "Scheduling support for concurrency and parallelism in the Mach operating system," *IEEE Computer*, vol. 23, no. 5, pp. 35–43, May 1990.
- [4] L. Kalita, "Socket Programming," *International Journal of Computer Science and Information Technologies*, vol. 5, no. 3, pp. 462–466, 2014.